

PRANAV SHARAN

Task 1 Prediction using Supervised ML

In this regression task we will predict the percentage of marks that a student is expected to score based upon the number of hours they studied.

We create a linear regression module to help us achieve this as the task involves only two variables.

Importing all required libraries

```
In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

Loading the dataset

```
In [3]: df = pd.read_csv("student_scores.csv")  
print("Data imported successfully")  
  
df.head(10)
```

Data imported successfully

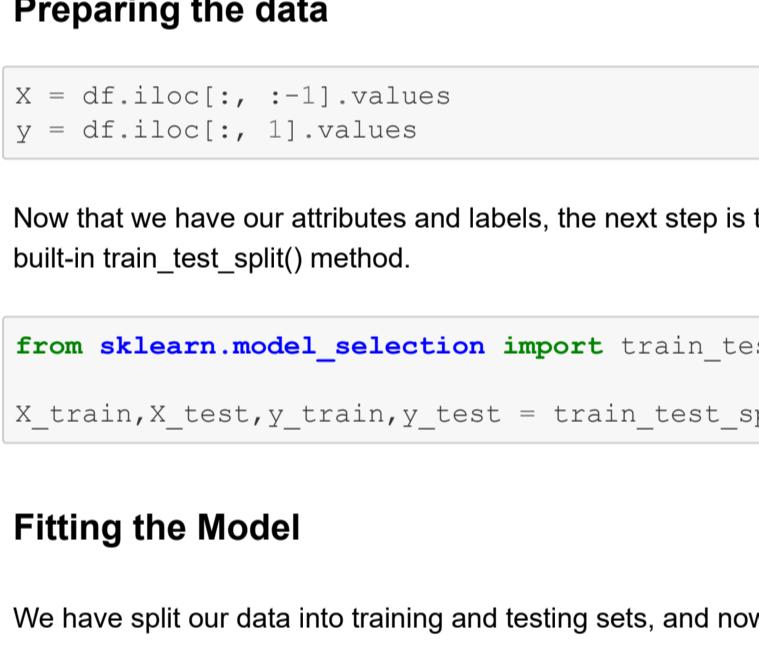
```
Out[3]:
```

Hours Scores

0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25

We plot our data points on 2-D graph to see if we can find any relationship between the Independent Variable and Dependent Variable. We use matplotlib to plot the graph.

```
In [4]: df.plot(x='Hours', y='Scores', style='o')  
plt.title('Hours vs Percentage')  
plt.xlabel('Hours Studied')  
plt.ylabel('Percentage Score')  
plt.show()
```



We can see that there is a positive linear relation between the number of hours studied and percentage of score.

Preparing the data

```
In [5]: X = df.iloc[:, :-1].values  
y = df.iloc[:, 1].values
```

Now that we have our attributes and labels, the next step is to split this data into training and test sets. We'll do this by using Scikit-Learn's built-in `train_test_split()` method.

```
In [6]: from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

Fitting the Model

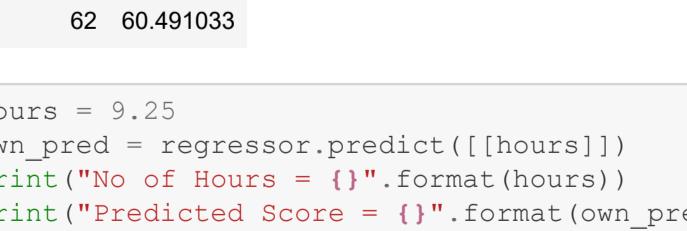
We have split our data into training and testing sets, and now is finally the time to train our algorithm.

```
In [8]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

```
Out[8]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [9]: # Plotting the regression line  
line = regressor.coef_*X+regressor.intercept_
```

```
# Plotting for the test data  
plt.scatter(X, y)  
plt.plot(X, line);  
plt.show()
```



Making Predictions

Now that we have trained our algorithm, it's time to make some predictions.

```
In [10]: y_pred = regressor.predict(X_test)  
df_pred = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
```

```
Out[10]:
```

Actual Predicted

0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
In [14]: hours = 9.25  
own_pred = regressor.predict([[hours]])  
print("No of Hours = {}".format(hours))  
print("Predicted Score = {}".format(own_pred[0]))
```

No of Hours = 9.25

Predicted Score = 93.69173248737538

Evaluating the model

```
In [15]: from sklearn import metrics  
print('Mean Absolute Error:',  
      metrics.mean_absolute_error(y_test, y_pred))
```

Mean Absolute Error: 4.183859899002975

```
In [ ]:
```